

A Novel Approach to Parameterized verification of Cache Coherence Protocols

Yongjian Li Kaiqiang Duan Shaowei Cai Yi Lv

Abstract—Parameterized verification of parameterized protocols like cache coherence protocols is important but hard. Our tool paraVerifier handles this hard problem in a unified framework: (1) it automatically discovers auxiliary invariants and the corresponding causal relations from a small reference instance of the verified protocol; (2) the above invariants and causal relation information are automatically generalized into a parameterized form to construct a parameterized formal proof in a theorem prover (e.g., Isabelle). The principle underlying the generalization is the symmetry mapping. Our method is successfully applied to typical benchmarks including snoopy-based and directory-based benchmarks. Another novel feature of our method lies in that the final verification result of a protocol is provided by a formal and readable proof.

I. INTRODUCTION

Verification of parameterized concurrent systems is interesting in the area of formal methods, mainly due to the practical importance of such systems. Parameterized systems exist in many important application areas, including cache coherence, security, and network communication protocols. The hardness of parameterized verification is mainly due to the requirement of correctness that the desired properties should hold in any instance of the parameterized system. The model checkers, although powerful in verification of non-parameterized systems, become impractical to verify parameterized systems, as they can verify only an instance of the parameterized system in each execution. A desirable approach is to provide a proof that the correctness holds for any instance.

Related Work: There have been a lot of studies in the field of parameterized verification [1], [2], [3], [4], [5], [6], [7], [8], [9]. Among them, the ‘invisible invariants’ method [3] is an automatic technique for parameterized verification. In this method, auxiliary invariants are computed in a finite system instance to aid inductive invariant checking. Combining parameter abstraction and guard strengthening with the idea of computing invariants in a finite instance, Lv et al. [8] use a small instance of a parameterized protocol as a ‘reference instance’ to compute candidate invariants. References to a specific node in these candidate invariants are then abstracted away, and the resulting formulas are used to strengthen guards of the transition rules in the abstract node. Both works [3], [8] attempt to automatically find invariants. However, the invisible invariants are raw boolean formulas transferred from the reachable state set of a small finite instance of a protocol, which are BDDs computed by TLV (an variant of BDD-based SMV model checker). They are too raw to have an intuitive meanings. The capacity of the invisible invariant method is seriously limited when computing the reachable set of invisible

invariants for the inductive checking is not feasible in the case of a large example like FLASH. Until now, the examples, which can be handled by the ‘invisible invariant’ method, are quite small, we still can’t find successful experiments on large examples like FLASH.

The CMP method, which adopts parameter abstraction and guard strengthening, is proposed in [6] for verifying a safety property *inv* of a parameterized system. An abstract instance of the parameterized protocol, which consists of $m + 1$ nodes $\{P_1, \dots, P_m, P^*\}$ with m normal nodes and one abstract node P^* , is constructed iteratively. The abstract system is an abstraction for any protocol instance whose size is greater than m . Normally the initial abstract system does not satisfy the invariant *inv*. Nevertheless it is still submitted to a model checker for verification. When a counterexample is produced, one needs to carefully analyze it and comes up with an auxiliary invariant *inv'*, then uses it to strengthen the guards of some transition rules of the abstract node. The ‘strengthened’ system is then subject to model checking again. This process stops until the refined abstract system eventually satisfies the original invariant as well as all the auxiliary invariants supplied by the user. However, this method’s soundness is only argued in an informal way. To the best of our knowledge, no one has formally proved its correctness in a theorem prover. This situation may be not ideal because its application domain for cache coherence protocols which demands the highest assurance for correctness. Besides, the analysis of counter-example and generation of new auxiliary invariants usually depend on human’s deep insightful understanding of the protocol. It is too laborious for people to do these analysis and some effective automatic tool is needed to help people.

Predicate abstraction is also applied to the verification of parameterized systems. Baukus, Lakhnech, and Stahl have used it to verify German (without data paths)[?], and Das, Dill, and Park have used it to verify FLASH[10]. The core of predicate abstraction is to discover a set of predicates, which are needed to abstract the states of a system, and an abstract state is a valuation of the predicates. Unfortunately, the task of discover proper predicates is neither easy nor automatic. Furthermore, the abstracted system is needed to proved to be conservative for certain properties under verification. This proof also needs a set of auxiliary invariants. Therefore searching enough auxiliary invariants can’t be avoided. No further efforts are made to make automatic both the discovery of proper predicates and the searching of auxiliary invariants in the work of applying predicate abstraction to the parameterized verification.

Sylvain Conchon, Amit Goel, Sava Krstic, Alain Mebsout, and Fatiha Zaidi have made progress in searching automatically auxiliary invariants[13]. A heuristics-guided algorithm, called Barb, searches auxiliary invariants backward with the help of an oracle (a reference instance of the protocol). Roughly speaking, Barb’s work can be seen as a backward reachability analysis. Barb is implemented in an SMT-based model checker Cubicle[?]. The correctness of Barb is argued in a generic symbolic framework. The searched auxiliary invariants are claimed to be inductive for deductive proof in a case study of a protocol. However, the formulation of Barb and the proof is not done in a theorem prover. Neither is a formal proof is given adopting the invariants for the protocol. Besides, the configuration of oracle need to be done manually.

The degree of scalability and automatic are two critical merits of approaches to parameterized verification. In this sense, verification of real-world parameterized systems is still a challenging task. For instance, up to now, the verification of a real-world benchmark FLASH requires human guidance in the existing successful verifications [10], [11], [6]. In order to effectively verify complex parameterized protocols like FLASH, there are two critical problems. The first one is how to find a set of sufficient and necessary invariants without (or with less) human intervention, which is a core problem in this field. The second one is the rigorousness of the verification. The theory foundation of a parameterized verification technique and its soundness are only discussed in a paper proof style in previous work. It is preferable to formulate all the verification in a publicly-recognized trust-worthy framework like a theorem prover [6]. However, theorem proving in a theorem prover like Isabelle is interactive, not automatical.

In order to solve the parameterized verification in a both automatical and rigorous way, we design a tool called **paraVerifier**, which is based on a simple but elegant theory. Three kinds of causal relations are introduced, which are essentially special cases of the general induction rule. Then, a so-called consistency lemma is proposed, which is the cornerstone in our method. Especially, the theory foundation itself is verified as a formal theory in Isabelle, which is the formal library for verifying protocol case studies. The library provides basic types and constant definitions to model protocol cases and lemmas to prove properties.

Our tool **paraVerifier** is composed of two parts: an invariant finder **invFinder** and a proof generator **proofGen**. Given a protocol \mathcal{P} and a property inv , **invFinder** tries to find useful auxiliary invariants and causal relations which are capable of proving inv . To construct auxiliary invariants and causal relations, we employ heuristics inspired by consistency relation. Also, when several candidate invariants are obtained using the heuristics, we use oracles such as a model checker and an SMT-solver to check each of them under a small reference model of \mathcal{P} , and chooses the one that has been verified.

After **invFinder** finds the auxiliary invariants and causal relations, **proofGen** generalizes them into a parameterized form, which are then used to construct a completely parameterized formal proof in a theorem prover (e.g., Isabelle) to

model \mathcal{P} and to prove the property inv . After the base theory is imported, the generated proof is checked automatically. Usually, a proof is done interactively. Special efforts in the design of the proof generation are made in order to make the proof checking automatically.

The organization of this work is as follows: Section II introduces the preliminaries; Section III introduces the theoretical foundation; Section IV the **invFinder**; Section V the generalization strategy; Section VI the **proofGen** and the generated proof. We go through these sections by verifying a small example - mutual exclusion example. Section VII shows the further experiments on real-world protocols. Section VIII concludes our work.

II. PRELIMINARIES

There are three kinds of *variables*: 1) simple identifier, denoted by a string; 2) element of an array, denoted by a string followed by a natural inside a square bracket. E.g., $arr[i]$ indicates the i th element of the array arr ; 3) filed of a record, denoted by a string followed by a dot and then another string. E.g., $red.f$ indicates the filed f of the record red . Each variable is associated with its *type*, which can be enumeration, natural number, and Boolean.

Expressions and *formulas* are defined mutually recursively. *Expressions* can be simple or compound. A simple expression is either a variable or a constant, while a compound expression is constructed with the *ite*(if-then-else) form $f ? e_1 : e_2$, where e_1 and e_2 are expressions, and f is a formula. A *formula* can be an atomic formula or a compound formula. An atomic formula can be a boolean variable or constant, or in the equivalence form $e_1 \doteq e_2$, where e_1 and e_2 are two expressions. A *formula* can also be constructed by using the logic connectives, including negation (\neg), conjunction (\wedge), disjunction (\vee), implication (\rightarrow).

An *assignment* is a mapping from a variable to an expression, and is denoted with the assigning operation symbol “ $:=$ ”. A *statement* α is a set of assignments which are executed in parallel, e.g., $x_1 := e_1; x_2 := e_2; \dots; x_k := e_k$. If an assignment maps a variable to a (constant) value, then we say it is a *value-assignment*. We use $\alpha|_x$ to denote the expression assigned to x under the statement α . For example, let α be $\{arr[1] := C; x := false\}$, then $\alpha|_x$ returns $false$. A *state* is an instantaneous snapshot of its behavior given by a set of value-assignments.

For every expression e and formula f , we denote the value of e (or f) under a state s as $\mathbb{A}[e, s]$ (or $\mathbb{B}[f, s]$). For a state s and a formula f , we write $s \models f$ to mean $\mathbb{B}[f, s] = true$. Formal semantics of expressions and formulas are given in HOL as usual, which is shown as follows: ¹

For an expression e and a statement $\alpha = x_1 := e_1; x_2 := e_2; \dots; x_k := e_k$, we use $\text{vars}(\alpha)$ to denote the variables to be assigned $\{x_1, x_2, \dots, x_k\}$; and use e^α to denote the

¹The logic to specify parameterized system can be embedded in HOL supported by Isabelle. Therefore, HOL can be regarded as the main meta-logic in our work.

Semantics
$\mathbb{A}[v, s] \equiv s(v)$, where v is a variable
$\mathbb{A}[c, s] \equiv c$, where c is a constant
$\mathbb{A}[f?e_1 : e_2, s] \equiv \text{if } (\mathbb{B}[f, s]) \text{ then } \mathbb{A}[e_1, s] \text{ else } \mathbb{A}[e_2, s]$
$\mathbb{B}[e_1 \doteq e_2, s] \equiv \mathbb{A}[e_1, s] = \mathbb{A}[e_2, s]$
$\mathbb{B}[\neg f, s] \equiv \neg \mathbb{B}[f, s]$
$\mathbb{B}[f_1 \wedge f_2, s] \equiv \mathbb{B}[f_1, s] \wedge \mathbb{B}[f_2, s]$
$\mathbb{B}[f_1 \vee f_2, s] \equiv \mathbb{B}[f_1, s] \vee \mathbb{B}[f_2, s]$
$\mathbb{B}[f_1 \rightarrow f_2, s] \equiv \mathbb{B}[f_1, s] \text{ implies } \mathbb{B}[f_2, s]$

```

assignN(i)  $\equiv$  n[i]=I
pini(N)  $\equiv$  x=true  $\wedge$  forallForm(N, assignN )
try(i)  $\equiv$  n[i]  $\doteq$  I  $\triangleright$  n[i] := T
crit(i)  $\equiv$  n[i]  $\doteq$  T  $\wedge$  x = true  $\triangleright$  n[i] := C; x := false
exit(i)  $\equiv$  n[i]  $\doteq$  C  $\triangleright$  n[i] := E
idle(i)  $\equiv$  n[i]  $\doteq$  E  $\triangleright$  n[i] := I; x := true
prules(N)  $\equiv$  {r.  $\exists i. i \leq N \wedge (r = \text{crit}(i) \vee r = \text{exit}(i))$ }
 $\vee r = \text{idle}(i)$   $\vee r = \text{try}(i)$ }
mutualEx(N)  $\equiv$  (pIni(N), prules(N))
mutualInv(i, j)  $\equiv$  ! (n[i]  $\doteq$  C  $\wedge$  n[j]  $\doteq$  C)

```

expression transformed from e by substituting each x_i with e_i simultaneously. Similarly, for a formula f and a statement $\alpha = x_1 := e_1; x_2 := e_2; \dots; x_k := e_k$, we use f^α to denote the formula transformed from f by substituting each x_i with e_i . Moreover, f^α can be regarded as the weakest precondition of formula f w.r.t. statement α , and we denote $\text{preCond}(f, \alpha) \equiv f^\alpha$. Noting that a state transition is caused by an execution of the statement, formally, we define: $s \xrightarrow{\alpha} s' \equiv (\forall x \in \text{vars}(\alpha). s'(x) = \mathbb{A}[\alpha|_x, s]) \wedge (\forall x \notin \text{vars}(\alpha). s'(x) = s(x))$.

A rule r is a pair $\langle g, \alpha \rangle$, where g is a formula and is called the guard of rule r , and α is a statement and is called the action of rule r . For convenience, we denote a rule with the guard g and the statement α as $g \triangleright \alpha$. Also, we denote $\text{act}(g \triangleright \alpha) \equiv \alpha$ and $\text{pre}(g \triangleright \alpha) \equiv g$. If the guard g is satisfied at state s , then α can be executed, thus a new state s' is derived, and we say the rule $g \triangleright \alpha$ is triggered at s , and transited to s' . Formally, we define: $s \xrightarrow{r} s' \equiv s \models \text{pre}(r) \wedge s \xrightarrow{\text{act}(r)} s'$.

A protocol \mathcal{P} is a pair (I, R) , where I is a set of formulas and is called the initializing formula set, and R is a set of rules. As usual, the reachable state set of protocol $\mathcal{P} = (I, R)$, denoted as $\text{reachableSet}(\mathcal{P})$, can be defined inductively: (1) a state s is in $\text{reachableSet}(\mathcal{P})$ if there exists a formula $f \in I$, and $s \models f$; (2) a state s is in $\text{reachableSet}(\mathcal{P})$ if there exists a state s_0 and a rule $r \in R$ such that $s_0 \in \text{reachableSet}(\mathcal{P})$ and $s_0 \xrightarrow{r} s$.

A parameterized object(T) is simple a function from a natural number to T, namely of type $\text{nat} \Rightarrow T$. For instance, a parameterized formula pf is of type $\text{nat} \Rightarrow \text{formula}$, and we define $\text{forallForm}(1, pf) \equiv pf(1)$, and $\text{forallForm}((n+1), pf) \equiv \text{forallForm}(n, pf) \wedge pf(n+1)$. $\text{existsForm}(1, pf) \equiv pf(1)$, and $\text{existsForm}((n+1), pf) \equiv \text{existsForm}(n, pf) \vee pf(n+1)$.

Now we use a simple example to illustrate the above definitions by a simple mutual exclusion protocol with N nodes. Let I, T, C, and E be three enumerating values, x, n are simple and array variables, N a natural number, $\text{pini}(N)$ the predicate to specify the initial state, $\text{prules}(N)$ the four rules of the protocol, $\text{mutualInv}(i, j)$ a property that $n[i]$ and $n[j]$ cannot be C at the same time. We want to verify that $\text{mutualInv}(i, j)$ holds for any $i \leq N, j \leq N$ s.t. $i \neq j$.

As Hoare logics specifies, after executing statement α , f holds iff $\text{preCond}(f, \alpha)$ holds before the execution.

Lemma 1 Suppose $s \xrightarrow{\alpha} s', s \models \text{preCond}(f, \alpha)$ if and only if $s' \models f$

III. CAUSAL RELATIONS AND CONSISTENCY LEMMA

A novel feature of our work lies in that three kinds of causal relations are exploited, which are essentially special cases of the general induction rule. Consider a rule r , a formula f , and a formula set fs , three kinds of causal relations are defined as follows:

Definition 1 We define the following relations:

$\text{invHoldRule}_1 :: \text{state} \times \text{formula} \times \text{rule} \Rightarrow \text{bool}$,
 $\text{invHoldRule}_2 :: \text{state} \times \text{formula} \times \text{rule} \Rightarrow \text{bool}$,
 $\text{invHoldRule}_3 :: \text{state} \times \text{formula} \times \text{rule} \times \text{ruleset} \Rightarrow \text{bool}$,
and $\text{invHoldRule} :: \text{state} \times \text{formula} \times \text{rule} \times \text{ruleset} \Rightarrow \text{bool}$.

- 1) $\text{invHoldRule}_1(s, f, r) \equiv s \models \text{pre}(r) \rightarrow s \models \text{preCond}(f, \text{act}(r));^2$
- 2) $\text{invHoldRule}_2(s, f, r) \equiv s \models f \leftrightarrow s \models \text{preCond}(f, \text{act}(r));$
- 3) $\text{invHoldRule}_3(s, f, r, fs) \equiv \exists f' \in fs \text{ s.t. } s \models (f' \wedge (\text{pre}(r) \rightarrow s \models \text{preCond}(f, \text{act}(r))))$
- 4) $\text{invHoldRule}(s, f, r, fs) \equiv s \models \text{invHoldRule}_1(s, f, r) \vee s \models \text{invHoldRule}_2(s, f, r) \vee s \models \text{invHoldRule}_3(s, f, r, fs)$.

The relation $\text{invHoldRule}(s, f, r, fs)$ defines a causality relation between f, r , and fs , which guarantees that if each formula in fs holds before the execution of rule r , then f holds after the execution of rule r . This includes three cases. 1) $\text{invHoldRule}_1(s, f, r)$ means that after rule r is executed, f becomes true immediately; 2) $\text{invHoldRule}_2(s, f, r)$ states that $\text{preCond}(S, f)$ is equivalent to f , which intuitively means that none of state variables in f is changed, and the execution of statement S does not affect the evaluation of f ; 3) $\text{invHoldRule}_3(s, f, r, fs)$ states that there exists another invariant $f' \in fs$ such that the conjunction of the guard of r and f' implies the precondition $\text{preCond}(S, f)$.

We can also view $\text{invHoldRule}(s, f, r, fs)$ as a special kind of inductive tactics, which can be applied to prove each formula in fs holds at each inductive protocol rule cases. Note that the three kinds of inductive tactics can be done by a theorem prover, which is the cornerstone of our work.

Example 1 Mutual-exclusion example.

²Here \rightarrow and \leftrightarrow are HOL connectives.

With the invHoldRule relation, we define a consistency relation $\text{consistent}(\text{invs}, \text{inis}, \text{rs})$ between a protocol (inis, rs) and a set of invariants $\text{invs} = \{\text{inv}_1, \dots, \text{inv}_n\}$.

Definition 2 We define a relation $\text{consistent} :: \text{formula set} \times \text{formula set} \times \text{rule set} \Rightarrow \text{bool}$. $\text{consistent}(\text{invs}, \text{inis}, \text{rs})$ holds if the following conditions hold:

- 1) for all formulas $\text{inv} \in \text{invs}$ and $\text{ini} \in \text{inis}$ and all states s , $s \models \text{ini}$ implies $s \models \text{inv}$;
- 2) for all formulas $\text{inv} \in \text{invs}$ and rules $r \in \text{rs}$ and all states s , $\text{invHoldRule}(s, \text{inv}, r, \text{invs})$

Example 2 Let us define a set of auxiliary invariants:

```

invOnXC(i)  $\equiv$   $!(x \doteq \text{true} \wedge n[i] \doteq C)$ 
invOnXE(i)  $\equiv$   $(x \doteq \text{true} \wedge n[i] \doteq E)$ 
aux1(i, j)  $\equiv$   $(n[i] \doteq C \wedge n[j] \doteq E)$ 
aux2(i, j)  $\equiv$   $(n[i] \doteq E \wedge n[j] \doteq E)$ 
pinvs(N)  $\equiv$   $\{f. \exists i \text{Inv1 } i \text{Inv2}. i \text{Inv1} \leq N \wedge i \text{Inv2} \leq N \wedge i \text{Inv1} \neq i \text{Inv2} \wedge f = \text{mutualInv } i \text{Inv1 } i \text{Inv2}\}$ 
 $\vee (\exists i \text{Inv1}. i \text{Inv1} \leq N \wedge f = \text{invOnXC } i \text{Inv1})$ 
 $\vee (\exists i \text{Inv1}. i \text{Inv1} \leq N \wedge f = \text{invOnXE } i \text{Inv1})$ 
 $\vee (\exists i \text{Inv1 } i \text{Inv2}. i \text{Inv1} \leq N \wedge i \text{Inv2} \leq N \wedge i \text{Inv1} \neq i \text{Inv2} \wedge f = \text{aux1 } i \text{Inv1 } i \text{Inv2})$ 
 $\vee (\exists i \text{Inv1 } i \text{Inv2}. i \text{Inv1} \leq N \wedge i \text{Inv2} \leq N \wedge i \text{Inv1} \neq i \text{Inv2} \wedge f = \text{aux2 } i \text{Inv1 } i \text{Inv2})$ 

```

In the following discussion, we assume that $\text{inv} = \text{mutual}(i_1, i_2)$, $r = \text{crit}(iR_1)$, $\text{rs} = \text{pinvs}(N)$, and assumptions $i_1 \neq N$, $i_2 \neq N$, $i_1 \neq i_2$, and $iR_1 \leq N$ hold.

$\text{noitemsep}, \text{no listsep}$

- $\text{invHoldRule}_2(s, \text{inv}, r)$, where $i_1 \neq iR_1$, and $i_2 \neq iR_1$, since $\text{preCond}(\text{act}(r), \text{inv}) = \text{inv}$.
- $\text{invHoldRule}_3(s, \text{inv}, r, \text{invs})$, where $i_1 = iR_1$. Since $\text{invOnXC}(i_2) \in \text{invs}$, $\text{preCond}(\text{act}(r), \text{inv}) = !(C \doteq C \wedge n[i_2] \doteq C)$, and $s \models (\text{invOnXC}(i_2) \wedge \text{pre}(\text{crit}(iR_1)))$ implies $s \models !(C \doteq C \wedge n[i_2] \doteq C)$.
- $\text{invHoldRule}_3(s, \text{inv}, r, \text{invs})$, where $i_2 = iR_1$. Since $\text{invOnXC}(i_1) \in \text{invs}$, $\text{preCond}(\text{act}(r), \text{inv}) = !(n[i_1] \doteq C \wedge C \doteq C)$, and $s \models (\text{invOnXC}(i_2) \wedge \text{pre}(\text{crit}(iR_1)))$ implies $s \models !(n[i_1] \doteq C \wedge C \doteq C)$.

For any invariant $\text{inv} \in \text{invs}$, inv holds at a reachable state s of a protocol $P = (\text{ini}, \text{rs})$ if the consistency relation $\text{consistent}(\text{invs}, \text{inis}, \text{rs})$ holds. The following lemma formalizes the essence of the aforementioned causal relation, and is called consistency lemma.

Lemma 2 If $P = (\text{ini}, \text{rs})$, $\text{consistent}(\text{invs}, \text{inis}, \text{rs})$, and $s \in \text{reachableSet}(P)$, then for all inv s.t. $\text{inv} \in \text{invs}$, $s \models \text{inv}$.

Now we apply the consistence lemma to prove that the mutual exclusion property holds for each reachable state of the mutual-exclusion protocol. Let us recall example 1 and 2.

Lemma 3 If $P = (\text{pini}(N), \text{prules}(N))$, $s \in \text{reachableSet}(P)$, and $0 < N$, then for all inv s.t. $\text{inv} \in \text{pinvs}(N)$, $s \models \text{inv}$.

Proof: By theorem2, we need to verify that parts (1) and (2) of the consistency relation hold.

$\text{noitemsep}, \text{no listsep}$

- (1) For any invariant $\text{inv} \in \text{pinvs}(N)$, any initializing predicate $\text{ini} \in \text{pini}(N)$, any state s , if ini is evaluated true at state s , then inv is evaluated true at state s .
- (2) For any invariant $\text{inv} \in \text{pinvs}(N)$, any r in rule set $\text{prules}(N)$, $\text{invHoldForRule}(s, \text{inv}, r, \text{pinvs}(N))$.

For (1), the proof is rather straightforward. For instance, consider the case where $\text{inv} = \text{mutualInv}(i_1, i_2)$ for some i_1 and i_2 , where $i_1 \leq N$, $i_2 \leq N$, and $i_1 \neq i_2$. We can conclude that $s \models n[i_2] \doteq I$ if $s \models \text{pini}(N)$, thus $s \models \text{inv}$ holds. The other invariants can be proved similarly.

For (2), we show proof of a typical case where $\text{inv} = \text{mutualInv}(i_1, i_2)$, and $r = \text{crit}(iR)$, where $iR \leq N$.

- (a) $iR = i_1$, we show $\text{invHoldForRule}_3(s, \text{inv}, r, \text{pinvs}(N))$. Because $\text{preCond}(\text{inv}, \text{act}(r)) = !(C \doteq C \wedge n[i_2] \doteq C)$ and $\text{pre}(r) = n[i_1] \doteq T \wedge x \doteq \text{ture}$, by IH, we have $s \models \text{invOnXC}(i_2)$, obviously, from this and $s \models \text{pre}(r)$, we have $s \models !(n[i_2] \doteq C)$, which is equivalent to $s \models \text{preCond}(\text{inv}, \text{act}(r))$.
- (b) $iR = i_2$, we show $\text{invHoldForRule}_3(s, \text{inv}, r, \text{pinvs}(N))$. Because $\text{preCond}(\text{inv}, \text{act}(r)) = !(n[i_1] \doteq C \wedge C \doteq C)$ and $\text{pre}(r) = n[i_2] \doteq T \wedge x \doteq \text{ture}$, by IH, we have $s \models \text{invOnXC}(i_1)$, obviously, from this and $s \models \text{pre}(r)$, we have $s \models !(n[i_1] \doteq C)$, which is equivalent to $s \models \text{preCond}(\text{inv}, \text{act}(r))$.
- (c) $iR \neq i_1$ and $iR \neq i_2$, we show $\text{invHoldForRule}_2(s, \text{inv}, r)$. The proof is straightforward because $\text{preCond}(\text{inv}, \text{act}(r)) = \text{inv}$. ■

In order to apply the consistency lemma to prove that a given property inv (e.g., the mutual exclusion property) holds for each reachable state of a protocol $P = (\text{inis}, \text{rs})$ (e.g., mutual-exclusion protocol), we need to solve two problems. First, we need to construct a set of auxiliary invariants invs which contains inv and satisfies $\text{consistent}(\text{invs}, \text{inis}, \text{rs})$. By applying the consistency lemma, we decompose the original problem of invariant checking into that of checking the causal relation between some $f \in \text{invs}$ and $r \in \text{rs}$. The latter needs three levels of case analysis: the first is on the form of inv , the second is on r , and the last is on the rule parameters iR and invariant parameters such as i_1 . Only if a proof script contains sufficient information on the case splitting and the kind of causal relation to be checked in each subcase, Isabelle can help us to automatically check it. How to generate automatically such a proof is the second problem.

Our solutions to the two problems are shown in Fig. 1: Given a protocol, invFinder finds all the necessary ground auxiliary invariants from a small instance of the




Fig. 1. The workflow of paraVerifier

protocol in Murphi. This step solves the first problem. A table `protocol.tbl` is worked out to store the set of ground invariants and causal relations, which are then used by `proofGen` to create an Isabelle proof script which models and verifies the protocol in a parameterized form. In this step, ground invariants are generalized into a parameterized form, and accordingly ground causal relations are adopted to create parameterized proof commands which essentially proves the existence of the parameterized causal relations. This solves the second problem. At last, the Isabelle proof script is fed into Isabelle to check the correctness of the protocol.

IV. SEARCHING AUXILIARY INVARIANTS

Given a protocol \mathcal{P} and a property set F containing invariant formulas we want to verify, `invFinder` in Algorithm 1 aims to find useful auxiliary invariants and causal relations which are capable of proving any element in F . A set A is used to store all the invariants found up to now, and is initialized as F . A queue $newInvs$ is used to store new invariants which have not been checked, and is initialized as F . A relation table $tuples$ is used to record the causal relation between a parameterized rule in some parameter setting and a concrete invariant. Initially $tuples$ is set as NULL. `invFinder` works iteratively in a semi-proving and semi-searching way. In each iteration, the head element f of $newInvs$ is popped, then `Policy`(r, f) generates groups of parameters $paras$ according to r and f by some policy. For each parameter $para$ in $paras$, it is applied to instantiate r into a concrete rule cr . Here `apply`($r, para$) = r if r contains no array-variables and $para = []$; otherwise `apply`($r, para$) = $r(para_{[1]}, \dots, para_{[|para|]})$. Then `coreFinder`(cr, f, A) is called to check whether a causal relation exists between cr and f ; if there is such one relation item, the relation item rel and a formula option $newInvOpt$ is returned; otherwise a run-time error occurs in `coreFinder`, which indicates no proof can be found. In the first case, a tuple $\langle r, para, f, rel \rangle$ will be inserted into $tuples$; If the formula option $newInvOpt$ is NONE, then no new invariant formula is generated; otherwise $newInvOpt = \text{Some}(f')$ for some formula f' , then `get`($newInvOpt$) returns f' , and the new invariant formula f' will be pushed into the queue $newInvs$ and inserted into the invariant set A . The above searching process is executed until $newInvs$ becomes empty. At last, the table $tuples$ is returned.

In Algorithm ??, the parameter generation policy `Policy` and the core invariant searching function `coreFinder` will be illustrated in Section IV-A and IV-B.

A. Parameter Generation Policy

In order to formulate our parameter generation policy, we introduce the concept of permutation modulo to symmetry relation \simeq_m^n , and a quotient set of perms_m^n (the set of all n -permutations of m) under the relation. Here an n -permutation

Algorithm 1: Algorithm: `invFinder`

Input: Initially given invariants F , a protocol $\mathcal{P} = \langle I, R \rangle$
Output: A set of tuples which represent causal relations between concrete rules and invariants:

```

1  $A \leftarrow F$ ;
2  $tuples \leftarrow []$ ;
3  $newInvs \leftarrow F$ ;
4 while  $newInvs$  is not empty do
5    $f \leftarrow newInvs.dequeue$ ;
6   for  $r \in R$  do
7      $paras \leftarrow \text{Policy}(r, f)$ ;
8     for  $para \in paras$  do
9        $cr \leftarrow \text{apply}(r, para)$ ;
10       $newInvOpt, rel \leftarrow \text{coreFinder}(cr, f, A)$ ;
11       $tuples \leftarrow tuples @ \langle r, para, f, rel \rangle$ ;
12      if  $newInvOpt \neq \text{NONE}$  then
13         $newInv \leftarrow \text{get}(newInvOpt)$ ;
14         $newInvs.enqueue(newInv)$ ;
15         $A \leftarrow A \cup \{newInv\}$ ;
16 return  $tuples$ ;

```

of m is ordered arrangement of an n -element subset of an m -element set $I = \{i.0 < i \leq m\}$. We use a list xs with size n to stand for a n -permutation of m . For instance, $[1, 2]$ is a 2-permutation of 3. $xs[i]$ and $|xs|$ denote the i -th element and the length of xs respectively. If $xs[i] = i$ for all $i \leq |xs|$, we call it identical permutation.

Definition 3 Let m and n be two natural numbers, where $n \leq m$, L and L' are two lists which stand for two n -permutations of m , $\text{noitemsep}, \text{noolistsep}$

- 1) $L \simeq_m^n L' \equiv (|L| = |L'| = n) \wedge (\forall i. i < |L| \wedge L[i] \leq m - n \rightarrow L[i] = L'[i])$.
- 2) $L \simeq_m^n L' \equiv L \simeq_m^n L' \wedge L' \simeq_m^n L$.
- 3) $\text{semiP}(m, n, S) \equiv (\forall L \in \text{perms}_m^n \exists L' \in S. L \simeq_m^n L') \wedge (\forall L \in S. \forall L' \in S. L \neq L' \rightarrow \neg(L \simeq_m^n L'))$.
- 4) A set S is called a quotient of the set perms_m^n under the relation \simeq_m^n if $\text{semiP}(m, n, S)$.

The definition of of relation \simeq_m^n (item 1 and 2 in Definition 3) directly leads to the following lemma.

Lemma 4 If $L \simeq_{m+n}^n L'$, then for any $0 < i \leq |L|$, any $0 < j \leq m$, $L[i] = j$ if and only if $L'[i] = j$.

For instance, let $L = [2, 3]$ and $L' = [2, 4]$, then $L \simeq_4^2 L'$. Due to Lemma 4, we can analyze a group of concrete parameters by analyzing only one of them as a presentative. Keeping this in mind, let us look at the following lemma, which together with Lemma 4 is the theoretical basis of our policy.

Lemma 5 Let S be a set s.t. $\text{semiP}(m, n, S)$, $\text{noitemsep}, \text{noolistsep}$

- 1) for any $L \in \text{perms}_m^n$, there exists a $L' \in S$ s.t. $L \simeq_m^n L'$.
- 2) let $L \in S, L' \in S$, if $L \neq L'$, then there exists two indice $i \leq m$ and $j \leq n$ such that $L[i] = j$ and $L'[i] \neq j$.

Lemma 5 shows 1) completeness of S w.r.t. the set perms_m^n under the relation \simeq , 2) the distinction between two different elements in S . Therefore, S has covered all analysing patterns according to the aforementioned comparing scheme between elements of L with numbers $j < n - m$. Moreover, the case patterns represented by different elements in S are different from each other. This fact can be illustrated by the following example.

Example 3 Let $m = 2$, $n = 1$, $S = \{[1], [2], [3]\}$ and $\text{semiP}(m, n, S)$, let LR be an element in S , there are three cases: *noitemsep, nolistsep*

- 1) $LR = [1]$: it is a special case where $LR_{[1]} = 1$;
- 2) $LR = [2]$: it is a special case where $LR_{[1]} = 2$;
- 3) $LR = [3]$: it is a special case where $LR_{[1]} \neq 1$ and $LR_{[1]} \neq 2$.

Note that the above cases are mutually disjoint, and their disjunction is true.

In Algorithm ??, a concrete formula cinv is popped from the queue newInvs , which can be seen as a normalized instantiation of some parameterized formula pinv .

Definition 4 A concrete invariant formula cinv is normalized w.r.t a parameterized invariant pinv if there exists no array variable in cinv and $\text{pinv} = \text{cinv}$ or there exists an identical permutation LI with $|LI| > 0$ such that $\text{cinv} = \text{pinv}(1, \dots, |LI|)$;

Any normalized cinv containing array variables is obtained by instantiating a parameterized invariant pinv with a parameter list which is an identical permutation LI (i.e., the j^{th} parameter is j itself $LI_{[j]} = j$). Thus, consider a list of parameter LR which is used to instantiate a parameterized rule pr , we have $LR_{[i]} = j$ (or $LR_{[i]} \neq j$) is equivalent to $LR_{[i]} = LI_{[j]}$ (or $LR_{[i]} \neq LI_{[j]}$), which is a factor to specify a case by comparing $LR_{[i]}$ with $LI_{[j]}$.

Let cinv be a normalized concrete invariant w.r.t. a parameterized invariant pinv , pr be a parameterized rule, m be the number of actual parameters occurring in cinv , and n be the number of formal parameters occurring in pr , our policy is to compute a quotient of perms_m^n , denoted as $\text{cmpSemiperm}(m + n, n)$, and use elements of it as a group of parameters to instantiate pr into a set crs of concrete rules.³ For instance, for the invariant $\text{mutuallnv}(1, 2)$, three groups of parameters $[1]$, $[2]$, $[3]$ are used to instantiate crit respectively, each of the instantiation results will be used to check which kind of causal relation exists between it and $\text{mutuallnv}(1, 2)$. Each of the three probed concrete causal relations will be used to generalized into a symbolic causal relation existing between crit and mutuallnv in a case formulated by a predicate comparing rule parameters and invariant parameters.

B. Core Searching Algorithm

For a cinv and a rule $r \in \text{crs}$, the core part of the invFinder tool is shown in Algorithm 2. It needs to call two oracles. The

first one, denoted by chk , checks whether a ground formula is an invariant. Such an oracle can be implemented by translating the formula into a formula in SMV, and calling SMV to check whether it is an invariant in a given small reference model of the protocol. If the reference model is too small to check the invariant, then the formula will be checked by Murphi in a big reference model. The second oracle, denoted by tautChk , checks whether a formula is a tautology. Such a tautology checker is implemented by translating the formula into a form in the SMT (SAT Modulo Theories) format, and checking it by an SMT solver such as Z3.

Algorithm 2: Core Searching Algorithm: *coreFinder*

Input: $r, \text{inv}, \text{invs}$
Output: A formula option f , a new causal relation rel

```

1  $g \leftarrow$  the guard of  $r$ ,  $S \leftarrow$  the statement of  $r$ ;
2  $\text{inv}' \leftarrow \text{preCond}(\text{inv}, S)$ ;
3 if  $\text{inv} = \text{inv}'$  then
4    $\text{relItem} \leftarrow (r, \text{inv}, \text{invRule}_2, -)$ ;
5   return (NONE,  $\text{relItem}$ );
6 else if  $\text{tautChk}(g \rightarrow \text{inv}') = \text{true}$  then
7    $\text{relItem} \leftarrow (r, \text{inv}, \text{invRule}_1, -)$ ;
8   return (NONE,  $\text{relItem}$ );
9 else
10   $\text{candidates} \leftarrow \text{subsets}(\text{decompose}(\text{dualNeg}(\text{inv}') \bar{\wedge} g))$ ;
11   $\text{newInv} \leftarrow \text{choose}(\text{chk}, \text{candidates})$ ;
12   $\text{relItem} \leftarrow (r, \text{inv}, \text{invRule}_3, \text{newInv})$ ;
13  if  $\text{isNew}(\text{newInv}, \text{invs})$  then
14     $\text{newInv} \leftarrow \text{normalize}(\text{newInv})$ ;
15    return (SOME( $\text{newInv}$ ),  $\text{relItem}$ );
16  else
17    return (NONE,  $\text{relItem}$ );
```

Input parameters of Algorithm 2 include a rule instance r , an invariant inv , a sets of invariants invs . The sets invs stores the auxiliary invariants constructed up to now. The algorithm searches for new invariants and constructs the causal relation between the rule instance r and the invariant inv . The algorithm returns a formula option and a causal relation item between r and inv . A formula option value NONE indicates that no new invariant is found, while SOME(f) indicates a new auxiliary invariant f is searched.

Algorithm *coreFinder* works as follows: after computing the pre-condition inv' (line 2), which is the weakest precondition of the input formula inv w.r.t. S , the algorithm takes further operations according to the cases it faces with:

- (1) If $\text{inv} = \text{inv}'$, meaning that statement S does not change inv , then no new invariant is created, and new causal relation item marked with tag invHoldRule_2 is recorded between r and inv .
- (2) If tautChk verifies that $g \rightarrow \text{inv}'$ is a tautology, then no new invariant is created, and the new causal relation item marked with tag invHoldRule_1 is recorded between r and inv .
- (3) If neither of the above two cases holds, then a new auxiliary invariant newInv will be constructed, which will make the causal relation invHoldRule_3 to hold. The candidate set is

³the details of computing $\text{cmpSemiperm}(m + n, n)$ can be found in [12].

$subsets(decompose(dualNeg(inv') \bar{\wedge} g))$, where $decompose(f)$ decompose f into a set of sub-formulas f_i such that each f_i is not of a conjunction form and f is semantically equivalent to $f_1 \bar{\wedge} f_2 \bar{\wedge} \dots \bar{\wedge} f_N$. $dualNeg(!f)$ returns f . $subsets(S)$ denotes the power set of S . A proper formula is chosen from the candidate set to construct a new invariant $newInv$. This is accomplished by the `choose` function, which calls the oracle `chk` to verify whether a formula is an invariant in the given reference model. After $newInv$ is chosen, the function `isNew` checks whether this invariant is new w.r.t. $newInvs$ or $invs$. If this is the case, the invariant will be normalized, and then be added into $newInvs$, and the new causal relation item marked with tag `invRule3` will be added into the causal relations. The meaning of the word “new” is modulo to the symmetry relation. E.g., $mutualInv(1, 2)$ is equivalent to $mutualInv(2, 1)$ in a symmetry view.

TABLE I
A FRAGMENT OF OUTPUT OF `invFinder`

rule	ruleParas	inv	causal relation	f'
..
crit	[1]	$mutualInv(1, 2)$	<code>invHoldRule3</code>	<code>invOnXC(2)</code>
crit	[2]	$mutualInv(1, 2)$	<code>invHoldRule3</code>	<code>invOnXC(1)</code>
crit	[3]	$mutualInv(1, 2)$	<code>invHoldRule2</code>	..
..
crit	[1]	<code>invOnXC(1)</code>	<code>invHoldRule1</code>	..
crit	[2]	<code>invOnXC(1)</code>	<code>invHoldRule1</code>	..

For instance, let $PR = \{try, crit, exit, idle\}$, $invs = \{mutualInv(1, 2)\}$, the output of the `invFinder`, which is stored in file `mutual.tbl`, is shown in Table I. In the table, each line records the index of a normalized invariant, name of a parameterized rule, the rule parameters to instantiate the rule, a causal relation between the ground invariant and a kind of causal relation which involves the kind and proper formulas f' in need (which are used to construct causal relations `invHoldRule3`). The auxiliary invariants found by `invFinder` include: $inv_2 \equiv !(x \doteq true \bar{\wedge} n[1] = C)$, $inv_3 \equiv !(n[1] = C \bar{\wedge} n[2] = E)$, $inv_4 \equiv !(x \doteq true \bar{\wedge} n[1] \doteq E)$, $inv_5 \equiv !(n[1] \doteq E \bar{\wedge} n[2] \doteq E)$.⁴

V. GENERALIZATION

Intuitively, generalization means that a concrete index (formula or rule) is generalized into a set of concrete indices (formulas or rules), which can be formalized by a symbolic index (formula or rules) with side conditions specified by constraint formulas. In order to do this, we adopt a new constructor to model symbolic index or symbolic value `symb(str)`, where str is a string. We use N to denote `symb("N")`, which formalizes the size of an parameterized protocol instance. A concrete index i can be transformed into a symbolic one by some special strategy g , namely $symbolize(g, i) = symb(g(i))$. In this work, two special transforming function $flnv(i) = "iInv" \hat{\text{itoea}}(i)$ and $flr(i) = "iR" \hat{\text{itoea}}(i)$, where $itoea(i)$ is

the standard function transforming an integer i into a string. We use special symbols $iInv_i$ to denote $symbolize(fInv, i)$; and iR_i to denote $symbolize(fIr, i)$. The former formalizes a symbolic parameter of a parameterized formula, and the latter a symbolic parameter of a parameterized rule. Accordingly, we define $symbolize2f(g, inv)$ (or $symbolize2r(g, r)$), which returns the symbolic transformation result to a concrete formula inv (or rule r) by replacing a concrete index i occurring in inv (or r) with a symbolic index $symbolize(g, i)$.

There are two main kinds of generalization in our work: (1) generalization of a normalized invariant into a symbolic one. The resulting symbolic invariants are used to create definitions of invariant formulas in Isabelle. For instance, $!(x \doteq true \bar{\wedge} n[1] \doteq C)$ is generalized into $!(x \doteq true \bar{\wedge} n[iInv_1] \doteq C)$. This kind of generalization is done with model constraints, which specify that any parameter index should be not greater than the instance size N , and parameters to instantiate a parameterized rule (formula) should be different. (2) The generalization of concrete causal relations into parameterized causal relations in Isabelle, and will be used in proofs of the existence of causal relations in Isabelle.

Since the first kind of generalization is simple, we focus on the second kind of generalization, which consists of two phases. Firstly, groups of rule parameters such as $[[1], [2], [3]]$ will be generalized into a list of symbolic formulas such as $[iR_1 \doteq iInv_1, iR_1 \doteq iInv_2, (iR_1 \neq iInv_1) \wedge (iR_1 \neq iInv_2)]$ ⁵, which stands for case-splittings by comparing a symbolic rule parameter iR_1 and invariant parameters $iInv_1$ and $iInv_2$. In the second phase, the formula field accompanied with a `invHoldRule3` relation is also generalized by some special strategy.

Now let us look at the first phase, starting with some definitions. Consider a line of concrete causal relation shown in Table I, there is a group of rule parameters LR , and a group of parameters LI occurring in an invariant formula.

Definition 5 Let LR be a permutation s.t. $|LR| > 0$, which represents a list of actual parameters to instantiate a parameterized rule, let LI be a permutation $|LI| > 0$, which represents a list of actual parameters to instantiate a parameterized invariant, we define: `noitemsep, nolistsep`

- 1) *symbolic comparison condition generalized from comparing $LR_{[i]}$ and $LI_{[j]}$:*
 $symbCmp(LR, LI, i, j) \equiv$

$$\begin{cases} iR_i \doteq iInv_j & \text{if } LR_{[i]} = LI_{[j]} \\ iR_i \neq iInv_j & \text{otherwise} \end{cases} \quad (1)$$

$$(2)$$

- 2) *symbolic comparison condition generalized from comparing $LR_{[i]}$ and with all $LI_{[j]}$:*
 $symbCaseL(LR, LI, i) \equiv$

$$\begin{cases} symbCmp(LR, LI, i, j) & \text{if } \exists! j. LR_{[i]} = LI_{[j]} \\ forallForm(|LI|, pf) & \text{otherwise} \end{cases} \quad (3)$$

$$(4)$$

where $pf(j) = symbCmp(LR, LI, i, j)$, and $\exists! j.P$ is an qualifier meaning that there exists a unique j s.t.

⁴The names `mutualEx` and `invOnX1` in this work are just for easy-reading, their index here is generated in some order by `invFinder`

⁵ $iR_1 \neq iInv_1$ is the abbreviation of $!(iR_1 \doteq iInv_1)$

property P ;

- 3) *symbolic case generalized from comparing LR with LI*
 $\text{symbCase}(LR, LI) \equiv \text{forallForm}(|LR|, pf)$, where
 $pf(i) = \text{symbCaseI}(LR, LI, i)$;
- 4) *symbolic partition generalized from comparing all $LRS_{[k]}$ with LI , where LRS is a list of permutations with the same length:*
 $\text{partition}(LRS, LI) \equiv \text{existsForm}(|LRS|, pf)$, where
 $pf(i) = \text{symbCase}(LRS_i, LI)$.

$\text{symbCmp}(LR, LI, i, j)$ defines a symbolic formula generalized from comparing $LR_{[i]}$ and $LI_{[j]}$; $\text{symbCaseI}(LR, LI, i)$ a symbolic formula summarizing the results of comparison between $LR_{[i]}$ and all $LI_{[j]}$ such that $j \leq |LI|$; $\text{symbCase}(LR, LI)$ a symbolic formula representing a subcase generalized from comparing all $LR_{[i]}$ and all $LI_{[j]}$; $\text{partition}(LRS, LI)$ is a disjunction of subcases $\text{symbCase}(LRS_{[i]}, LI)$. Recall the first three lines in Table I, and $LI = [1, 2]$ is the list of parameters occurring in $\text{mutualEx}(1, 2)$; and LR is the actual parameter list to instantiate crit .

$\text{noitemsep, nolistsep}$

- when $LR = [1]$, $\text{symbCmp}(LR, LI, 1, 1) = (iR_1 \doteq iInv_1)$, $\text{symbCase}(LR, LI) = \text{symbCaseI}(LR, LI, 1) = (iR_1 \doteq iInv_1)$ because $LR_{[1]} = LI_{[1]}$.
- when $LR = [2]$, $\text{symbCmp}(LR, LI, 1, 2) = (iR_1 \doteq iInv_2)$, $\text{symbCase}(LR, LI) = \text{symbCaseI}(LR, LI, 2) = (iR_1 \doteq iInv_2)$ because $LR_{[1]} = LI_{[2]}$.
- when $LR = [3]$, $\text{symbCmp}(LR, LI, 1, 1) = (iR_1 \neq iInv_1)$, $\text{symbCmp}(LR, LI, 1, 2) = (iR_1 \neq iInv_2)$, $\text{symbCase}(LR, LI) = \text{symbCaseI}(LR, LI, 1) = (iR_1 \neq iInv_1) \wedge (iR_1 \neq iInv_2)$ because neither $LR_{[1]} = LI_{[1]}$ nor $LR_{[1]} = LI_{[2]}$.
- let $LRS = [[1], [2], [3]]$, $\text{partition}(LRS, LI) = (iR_1 \doteq iInv_1) \vee (iR_1 \doteq iInv_2) \vee ((iR_1 \neq iInv_1) \wedge (iR_1 \neq iInv_2))$

If we see a line in table I as a concrete test case for some concrete causal relation, then $\text{symbCase}(LR, LI)$ is an abstraction predicate to generalize the concrete case. Namely, if we transform $\text{symbCase}(LR, LI)$ by substituting $iInv_i$ with $LI_{[i]}$, and iR_j with $LR_{[j]}$, the result is semantically equivalent to true.

The second phase of generalization of concrete causal relations is to generalize the formula inv' accompanied with a causal relation invHoldRule_3 in a line of table I. An index occurring in f' can either occur in the invariant formula, or in the rule. We need to look it up to determine the transformation.

Definition 6 Let LI and LR are two permutations, $\text{find_first}(L, i)$ returns the least index j s.t. $L_{[j]} = i$ if there exists such an index; otherwise returns an error.

$$\text{lookup}(LI, LR, i) \equiv \begin{cases} iInv_{\text{find_first}(LI, i)} & \text{if } i \in LI \\ iR_{\text{find_first}(LR, i)} & \text{otherwise} \end{cases} \quad (6)$$

$\text{lookup}(LI, LR, i)$ returns the symbolic index transformed from i according to whether i occurs in LI or in LR .

The index i will be transformed into $iInv_{\text{find_first}(LI, i)}$ if i occurs in LI , and $iR_{\text{find_first}(LR, i)}$ otherwise. Employing the lookup strategy to transform a concrete index i in inv' to $\text{lookup}(LI, LR, i)$, symbolize2f transforms inv' into a symbolic one which will be needed in a proof command for existence of the invHoldRule_3 relation in Isabelle.

VI. AUTOMATICAL GENERATION OF ISABELLE PROOF

A formal model for a protocol case in a theorem prover like Isabelle includes the definitions of constants and rules and invariants, lemmas, and proofs. Readers can refer to [12] for detailed illustration of the formal proof script. In this section, we focus on the generation of a lemma on the existence of causal relation between a parameterize rule and invariant formula based on the aforementioned generalization of lines of concrete causal relations.

An example lemma critVsinv_1 and its proof in Isabelle in the mutualEx protocol, is illustrated as follows:

```
1 lemma critVsinv1:
2   assumes a1:  $\exists iR1. iR1 \leq N \wedge r = \text{crit } iR1$  and
3   a2:  $\exists iInv1 \ iInv2. iInv1 \leq N \wedge iInv2 \leq N \wedge iInv1 \neq iInv2$ 
4    $\wedge f = \text{inv1 } iInv1 \ iInv2$ 
5   shows  $\text{invHoldRule } s \ f \ r \ (\text{invariants } N)$ 
6 proof -
7   from a1 obtain iR1 where a1:  $iR1 \leq N \wedge r = \text{crit } iR1$ 
8   by blast
9   from a2 obtain iInv1 iInv2 where a2:  $iInv1 \leq N$ 
10   $\wedge iInv2 \leq N \wedge iInv1 \neq iInv2 \wedge f = \text{inv1 } iInv1 \ iInv2$ 
11  by blast
12 5 have  $iR1 = iInv1 \vee iR1 = iInv2 \vee (iR1 \neq iInv1 \wedge iR1 \neq iInv2)$ 
13  by auto
14 6 moreover {assume b1:  $iR1 = iInv1$ 
15    7   have  $\text{invHoldRule3 } s \ f \ r \ (\text{invariants } N)$ 
16    proof (cut_tac a1 a2 b1, simp,
17      rule_tac x =! (x = true  $\wedge$  n[iInv2] = C) in exI, auto) qed
18    8   then have  $\text{invHoldRule } s \ f \ r \ (\text{invariants } N)$  by auto
19    9 moreover {assume b1:  $iR1 = iInv2$ 
20      10   have  $\text{invHoldRule3 } s \ f \ r \ (\text{invariants } N)$ 
21      proof (cut_tac a1 a2 b1, simp,
22        rule_tac x =! (x = true  $\wedge$  n[iInv1] = C) in exI, auto) qed
23      11   then have  $\text{invHoldRule } s \ f \ r \ (\text{invariants } N)$  by auto
24      12 moreover {assume b1:  $(iR1 \neq iInv1 \wedge iR1 \neq iInv2)$ 
25        13   have  $\text{invHoldRule2 } s \ f \ r$ 
26        proof (cut_tac a1 a2 b1, auto) qed
27        14   then have  $\text{invHoldRule } s \ f \ r \ (\text{invariants } N)$  by auto
28        15 ultimately show  $\text{invHoldRule } s \ f \ r \ (\text{invariants } N)$  by blast
29        16 qed
```

In the above proof, line 2 are assumptions on the parameters of the invariant and rule, which are composed of two parts: (1) assumption $a1$ specifies that there exists an actual parameter $iR1$ with which r is a rule obtained by instantiating crit ; (2) assumption $a2$ specifies that there exists actual parameters $iInv1$ and $iInv2$ with which f is a formula obtained by instantiating inv1 . Line 4 are two typical proof patterns forward-style which fixes local variables such as $iR1$ and new facts such as $a1: iR1 \leq N \wedge r = \text{crit } iR1$. From line 5, the remaining part is a typically readable Isar proof using calculation reasoning such as moreover and ultimately to do case analysis. Line 5 splits cases of $iR1$ into all possible cases by comparing $iR1$ with $iInv1$ and $iInv2$, which is in fact characterized by $\text{partition}([1], [2], [3]), [1, 2])$. Lines 6-14 proves these cases one by one: Lines 6-8 proves the case where $iR1 = iInv1$, line 7 first proves that the causal relation invHoldRule_3 holds by supplying a symbolic formula, which is transformed from $\text{invOnXC}(2)$ by calling symbolize2f

with lookUp strategy. From the conclusion at line 7, line 8 furthermore proves the causal relation `invHoldRule` holds; Lines 9-11 proves the case where `iR1=iInv2`, proof of which is similar to that of case 1; Lines 12-14 the case where neither `iR1=iInv1` nor `iR1=iInv2`. Each proof of a subcase is done in a block moreover `b1:asml proof1`, the ultimately proof command in line 15 concludes by summing up all the subcases.

With the help of all the lemmas such as `ruleVsinvl`, we can prove the following lemma `lemma_inv_1_on_rules` which specifies that for all $r \in \text{rules } N$, and f is a formula f which is generated by instantiating `inv1` with some parameters $iInv_1$ and $iInv_2$, `invHoldForRule s f r (invariants N)`.

```
lemma lemma_inv1_on_rules:
  assumes a1: r ∈ rules N and
  a2: (∃ _iInv1 _iInv2. _iInv1 ≤ N ∧ _iInv2 ≤
    N ∧ iInv1 ≠ iInv2 ∧ f = inv1 iInv1 iInv2)
  shows invHoldForRule s f r (invariants N)
  proof -
    have (∃ i. i ≤ N ∧ r = try i) ∨ (∃ i. i ≤
      N ∧ r = crit i) ∨ (∃ i. i ≤ N ∧ r = exit i) ∨
      (∃ i. i ≤ N ∧ r = idle i)
    apply (cut_tac a1, auto) done
  moreover { assume b1: (∃ i. i ≤ N ∧ r = try i)
    have invHoldForRule' s f r (invariants N)
    apply (cut_tac a2 b1, metis tryVsinvl) done }
  moreover { assume a1: (∃ i. i ≤ N ∧ r = crit i)
    have invHoldForRule' s f r (invariants N)
    apply (cut_tac a2 b1, metis critVsinvl) done }
  moreover { assume a1: (∃ i. i ≤ N ∧ r = exit i)
    have invHoldForRule' s f r (invariants N)
    apply (cut_tac a2 b1, metis exitVsinvl) done }
  moreover { assume a1: (∃ i. i ≤ N ∧ r = idle i)
    have invHoldForRule' s f r (invariants N)
    apply (cut_tac a2 b1, metis idleVsinvl) done }
  ultimately show invHoldForRule' s f r (invariants N)
  by auto
qed
```

With the help of all the lemmas such as `lemma_invi_on_rules`, we can prove the following lemma `invs_on_rules` which specifies that for all $f \in \text{invariants } N$ and $r \in \text{rules } N$, `invHoldForRule s f r (invariants N)`.

```
lemma invs_on_rules: assumes a1: f ∈ invariants N
  and a2: r ∈ rules N
  shows invHoldForRule' s f r (invariants N)
  proof -
    have b1: (∃ iInv1 iInv2. iInv1 ≤ N ∧ iInv2 ≤ N ∧ iInv1 ≠ iInv2 ∧
      f = inv1 iInv1 iInv2) ∨
      (∃ iInv2. iInv2 ≤ N ∧ f = inv2 iInv2) ∨
      (∃ iInv1 iInv2. iInv1 ≤ N ∧ iInv2 ≤ N ∧ iInv1 ≠ iInv2 ∧
      f = inv3 iInv1 iInv2) ∨
      (∃ iInv2. iInv2 ≤ N ∧ f = inv4 iInv2) ∨
      (∃ iInv1 iInv2. iInv1 ≤ N ∧ iInv2 ≤ N ∧ iInv1 ≠ iInv2 ∧
      f = inv5 iInv1 iInv2)
    apply (cut_tac a1, auto) done
  moreover { assume b1: (∃ iInv1 iInv2. iInv1 ≤ N ∧ iInv2 ≤ N
    ∧ iInv1 ≠ iInv2 ∧ f = inv1 iInv1 iInv2)
    have invHoldForRule' s f r (invariants N)
    apply (cut_tac a2 b1, metis lemma_inv1_on_rules) done }
  moreover { assume b1: (∃ iInv2. iInv2 ≤ N
    ∧ f = inv2 iInv2)
    have invHoldForRule' s f r (invariants N)
    apply (cut_tac a2 b1, metis lemma_inv2_on_rules) done }
  moreover { assume b1: (∃ iInv1 iInv2. iInv1 ≤ N ∧ iInv2 ≤ N
    ∧ iInv1 ≠ iInv2 ∧ f = inv3 iInv1 iInv2)
    have invHoldForRule' s f r (invariants N)
    apply (cut_tac a2 b1, metis lemma_inv3_on_rules) done }
  moreover { assume b1: (∃ iInv2. iInv2 ≤ N ∧ f = inv4 iInv2)
    have invHoldForRule' s f r (invariants N)
    apply (cut_tac a2 b1, metis lemma_inv4_on_rules) done }
  moreover { assume b1: (∃ iInv1 iInv2. iInv1 ≤ N ∧ iInv2 ≤ N
    ∧ iInv1 ≠ iInv2 ∧ f = inv5 iInv1 iInv2)
    have invHoldForRule' s f r (invariants N)
    apply (cut_tac a2 b1, metis lemma_inv5_on_rules) done }
  ultimately show invHoldForRule' s f r (invariants N)
  apply fastforce done
qed end
```

1) Lemmas on initial states: In this section, we discuss the definition on the initial state of the protocol, and the lemmas specifying that each invariant formula holds at the initial state.

A typical Isabelle definition on the initial state of the protocol is as follows:

```
definition initState0::nat ⇒ formula where [simp]:
  initState0 N ≡ (forallForm (down N)
    (% i . (eqn (IVar (Para (Ident 'n')) i)) (Const I))))
definition initState1::formula where [simp]:
  initState1 ≡ (eqn (IVar (Ident 'x')) (Const true))
definition allInitSpecs::nat Rightarrow formula list
  allInitSpecs N ≡ [(initSpec0 N), (initSpec1)]
lemma iniImPLY_inv4:
  assumes a1: (∃ iInv1. iInv1 ≤ N ∧ f = inv4 iInv1)
  and a2: formEval (andList (allInitSpecs N)) s
  shows formEval f s
  using a1 a2 by auto
```

`initSpec0` and `initSpec1` specifies the assignments on each variable $n[i]$ where $i \leq N$ and x . The specifications of the initial state is the list of all the specification definition on related state variables. Lemma `iniImPLY_inv4` simply specifies that the invariant formula `inv4` holds at a state s which satisfies the conjunction of the specification of the initial state. Isabelle's `auto` method can solve this goal automatically. Other lemmas specifying that other invariant formulas hold at the initial state are similar.

With the lemmas such as `iniImPLY_inv4`, for any invariant $inv \in (\text{invariants } N)$, any state s , if `ini` is evaluated true at state s , then `inv` is evaluated true at state s .

```

lemma on_inis: assumes a1: f ∈ (invariants N)
and a2: ini ∈ { andList (allInitSpecs N) }
and a3: formEval ini s
shows formEval f s
proof -
  have c1: (∃ iInv1 iInv2. iInv1 ≤ NiInv2 ≤ NiInv1 ≠ iInv2
    ∧ f = inv__1 iInv1 iInv2) ∨
    (∃ iInv2. iInv2 ≤ N ∧ f = inv__2 iInv2) ∨
    (∃ iInv1 iInv2. iInv1 ≤ NiInv2 ≤ NiInv1 ≠ iInv2
    ∧ f = inv__3 iInv1 iInv2) ∨
    (∃ iInv2. iInv2 ≤ N ∧ f = inv__4 iInv2) ∨
    (∃ iInv1 iInv2. iInv1 ≤ NiInv2 ≤ NiInv1 ≠ iInv2
    ∧ f = inv__5 iInv1 iInv2)
  apply (cut_tac a1, simp) done
  moreover { assume b1: (∃ iInv1 iInv2. iInv1 ≤ NiInv2 ≤ N
    ∧ iInv1 ≠ iInv2 ∧ f = inv__1 iInv1 iInv2)
    have formEval f s
      apply (rule iniImpl_inv__1)
      apply (cut_tac b1, assumption)
      apply (cut_tac a2 a3, blast) done }
  moreover { assume b1: (∃ iInv2. iInv2 ≤ N ∧ f = inv__2 iInv2)
    have formEval f s
      apply (rule iniImpl_inv__2)
      apply (cut_tac b1, assumption)
      apply (cut_tac a2 a3, blast) done }
  }
  moreover { assume b1: (∃ iInv1 iInv2. iInv1 ≤ NiInv2 ≤ N
    ∧ iInv1 ≠ iInv2 ∧ f = inv__3 iInv1 iInv2)
    have formEval f s
      apply (rule iniImpl_inv__3)
      apply (cut_tac b1, assumption)
      apply (cut_tac a2 a3, blast) done }
  moreover { assume b1: (∃ iInv2. iInv2 ≤ N ∧ f = inv__4 iInv2)
    have formEval f s
      apply (rule iniImpl_inv__4)
      apply (cut_tac b1, assumption)
      apply (cut_tac a2 a3, blast) done }
  moreover { assume b1: (∃ iInv1 iInv2. iInv1 ≤ NiInv2 ≤ N
    ∧ iInv1 ≠ iInv2 ∧ f = inv__5 iInv1 iInv2)
    have formEval f s
      apply (rule iniImpl_inv__5)
      apply (cut_tac b1, assumption)
      apply (cut_tac a2 a3, blast) done }
  ultimately show formEval f s by auto
qed

```

The proof structure of lemma `inv1_on_rules` and `invs_on_rules` and `on_inis` are also typical case analysis ones using `moreover` blocks and `ultimately` commands, therefore, a generic program of generating a typical case analysis proof will be adopted in our framework.

2) *The main theorem:* With the preparation of lemma `on_inis` and `invs_on_rules`, the generation of the main lemma is quite easy. Recall that the consistency lemma is our main weapon to prove the main lemma, which requires proving two parts of obligations.

- (1) For any invariant $inv \in (\text{invariants } N)$, any state s , if ini is evaluated true at state s , then inv is evaluated true at state s . This can be solved done by applying lemma `on_inis`.
- (2) For any invariant $inv \in (\text{invariants } N)$, any r in rule set $rules\ N$, one of the causal relations $invHoldForRule_{1-3}$ holds. This can be solved done by applying lemma `invs_on_rules`.

```

lemma main: assumes a1: 0 < N and
a2: s ∈ reachableSet {andList (allInitSpecs N)} (rules N)
shows ∀ inv. inv ∈ (invariants N) → formEval inv s
proof (rule consistentLemma)
  show consistent (invariants N) {andList (allInitSpecs N)}
  (rules N)
  proof (cut_tac a1, unfold consistent_def, rule conjI)
    show ∀ inv ini s. inv ∈ (invariants N) → ini ∈ {andList
      (allInitSpecs N)} → formEval ini s → formEval inv s
    proof ((rule allI)+, (rule impI)+)
      fix inv ini s
      assume b1: inv ∈ (invariants N) and b2: formEval ini s
      and b3: ini ∈ {andList (allInitSpecs N)}
      show "formEval f s"
      apply (rule on_inis, cut_tac b1, assumption, cut_tac b2,
        assumption, cut_tac b3, assumption) done
    qed
  next show ∀ inv r. inv ∈ invariants N → r ∈ rules N
    → invHoldForRule inv r (invariants N)
    proof ((rule allI)+, (rule impI)+)
      fix f r
      assume b1: f ∈ invariants N and b2: r ∈ rules N
      show invHoldForRule' s f r (invariants N)
      apply (rule invs_on_rules, cut_tac b1, assumption,
        cut_tac b2, assumption) done
    qed
  next show s ∈ reachableSet andList (allInitSpecs N) (rules N)
    apply (metis a1) done
  qed

```

The generation of the main lemma is quite easy because it is in a standard form.

A. Algorithms of Proof Generator `proofGen`

In this subsection, we illustrate the key techniques and algorithms of generation of the lemmas and their proofs in subsection ?? . Being according with the order in which we introduce the above lemmas, we also introduce their generation in a bottom-up order. First let us introduce the generation of a subproof according to a relation tag of $invHoldForRule_{1-3}$, which is shown in Algorithm 10.

Algorithm 3: Generating a kind of proof which is according with a relation tag of $invHoldForRule_{1-3} : rel2proof$

Input: A causal relation item $relTag$

Output: An Isabelle proof: $proof$

```

1 if  $relTag = invHoldForRule_1$  then
2    $proof \leftarrow \text{sprintf}$ 
3     "have invHoldForRule1 f r (invariants N)
4     by(cut_tac a1 a2 b1, simp, auto)
5     then have invHoldForRule f r (invariants N) by blast" ;
6 else if  $relTag = invHoldForRule_2$  then
7    $proof \leftarrow \text{sprintf}$ 
8     "have invHoldForRule2 f r (invariants N) by(cut_tac a1
9     a2 b1, simp, auto)
10    then have invHoldForRule f r (invariants N) by blast" ;
11 else
12    $f' \leftarrow getFormField(relTag)$ ;
13    $proof \leftarrow \text{sprintf}$ 
14     "have invHoldForRule3 f r (invariants N)
15     proof(cut_tac a1 a2 b1, simp, rule_tac x=%s in
16     exI, auto)qed
17     then have invHoldForRule f r (invariants N) by blast"
18     (symbf2Isabelle f)";
19 return  $proof$ 

```

In the body of function `rel2proof`, `sprintf` writes a formatted data to string and returns it. In line

10, *getFormField(relTag)* returns f' if $relTag = invHoldForRule_3(f')$. *rel2proof* transforms a relation tag into a paragraph of proof. If the tag is among $invHoldForRule_{1-2}$, the transformation is rather straight-forward, else the form f' is assigned by the formula *getFormField(relTag)*, and provided to tell Isabelle the formula which should be used to construct the *invHoldForRule₃* relation.

Algorithm 4: Generating one sub-proof for a subcase: *oneMoreOverGen*

Input: A formula *caseFsm* standing for the assumption of the subcase, a relation item *relItem* containing the information of causal relation

Output: An Isabelle proof: *subProof*

```

1 proof ← rel2proof(relItem);
2 subProof ← sprintf
3   "moreover{assume b1:%s
4   %s } "
5   (asm, proof);
6 return subproof

```

In Algorithm 4, *oneMoreOverGen* generates a subproof for a subcase in a proof of case analysis. It returns a subproof which is composed by filling an assumption of the subcase such as "iR1=iInv1" and a paragraph of proof generated by *rel2proof(relItem)* into a format of block *moreover { ... }*.

Due to the common use of case analysis proof of using *moreover* and ultimately commands, we design a generic program of generating doing case analysis *doCaseAnalz*. In algorithm 5, formulas standing for case-splitting *partition*, subproofs *subproofs*, and the conclusion *concluding* are needed in case analysis to fill the format.

Algorithm 5: Generating a whole proof of doing case analysis: *doCaseAnalz*

Input: A formula *partition* standing for case-splittings, a proof list *subproofs* standing all the subproofs of each subcases, concluding parts *concluding*

Output: An Isabelle proof: *proof*

```

1 proof ← sprintf
2   " have %s by auto
3   %s
4   ultimately show %s by auto"
5   (partition, subproofs, concluding) ;
6 return proof

```

In algorithm 6, *caseAnalzI* generates a typical proof of doing case analysis to prove some causal relation hold between some rule and invariant. *oneMoreOverGenI(case,rel)* formula comes from the disjunction of formulas in the *symbCases* field of *rec*, which is returned by *caseField(rec)*, subproofs *subproofs* are generated by concatenation of all the subproofs, each of which is generated by *oneMoreOverGenI(case,rel)*. The proof is simply composed by calling *doCaseAnalz(partition, subproofs, concluding)*.

Algorithm 6: Generating a whole proof of doing case analysis on parameters of rule and invariant: *caseAnalzI*

Input: A record *rec* fetched from *symbCausal*

Output: An Isabelle proof: *proof*

```

1 cases ← caseField(rec);
2 rels ← relItems(rec); partition ← ∨ cases;
3 subproofs ← "";
4 while (cases ≠ []) do
5   case ← hd(cases) ;
6   cases ← tl(cases) ;
7   rel ← hd(rels) ;
8   rels ← tl(rels) ;
9   subproofs ←
    subproofs ^ oneMoreOverGenI(case, rel);
10 concluding ← "invHoldForRule s f r (invariants N) ";
11 proof ← doCaseAnalz(partition, subproofs, concluding);
12 return proof

```

Next we discuss how to generate assumptions on an invariant formula of an lemma such as *critVsInv1*. In the body of algorithm 7, *tbl_element(symbInvs, invName)* retrieves the record on a invariant formula from *symbInvs* to *invItem* by its name *invName*, *invParaNum(invItem)* and *constrOfInv(invItem)* return the field *invNumFld* and *constr* of *invItem* respectively. *invParasGen(lenPIinv)* generates a string of a list of actual parameters such as *iInv1...iInvlenPIinv* if *lenPIinv* > 0, else an empty string "". At last, the assumption on the invariant is created by filling *invParas*, *constrOnInv*, and *invName* into a proper place in the format if needed.

Algorithm 7: Generating an assumption on an invariant formula: *asmGenOnInv*

Input: An invariant name *invName*, a table *symInvs* storing invariant formulas

Output: An assumption on an invariant formula: *asm*

```

1 invItem ← tbl_element(symbInvs, invName);
2 lenPIinv ← invParaNum(invItem);
3 invParas ← invParasGen(lenPIinv);
4 constrOnInv ←
  symbForm2Isabelle(constrOfInv(invItem));
5 if lenPIinv = 0 then
6   asm ← "a1 : f = " ^ invName;
7 else
8   asm ← sprintf "a1 : ∃ %s. %s ∧ f=%s %s" (invParas,
    constrOnInv, invName, invParas);
9 return asm

```

Similar to *asmGenOnInv*, *obtainGenOnInv*, which is shown in algorithm 8, generates a proof command of *obtain* by retrieving and generating the related information and filling them in a format on *obtain*. Similar to *asmGenOnInv* and *obtainGenOnInv*, *asmGenOnRule* and *obtainGenOnRule* generate an assumption and *obtain* proof command on a rule.

After the above preparing functions, now the generation of a lemma on the causal relation such as *critVsInv1* is rather easy, which is shown in algorithm 9. After generating an assumption on invariant formula *asm1*, *asm2* on a rule, an *obtain* command *obtain1* on the invariant, and *obtain2*

Algorithm 8: Generating an obtain proof command on an invariant formula: obtainGenOnInv

Input: An invariant name *invName*, a table *symInvs* storing rules

```

1 invItem  $\leftarrow$  tbl_element(symInvs, invName);
2 lenPInv  $\leftarrow$  invParaNum(invItem);
3 invParas  $\leftarrow$  invParasGen(lenPInv);
4 if lenPInv = 0 then
5   | obtain  $\leftarrow$  "";
6 else
7   | obtain  $\leftarrow$  sprintf "from a1 obtain %s where a1:%s  $\wedge$  f=%s
   | %s by auto"
8   | (invParas, constrOnInv, invName, invParas);
9 return obtain

```

on the rule, *symRelItem* is retrieved from *symCausalTab* by *ruleName* \wedge *invName*, and a proof *proof* is generated by calling *caseAnalzI*(*symRelItem*). At last these parts are filled into proper places in the lemma format.

Algorithm 9: Generating a lemma on a causal relation: lemmaOnCausalRuleInv

Input: A parameterized rule name *ruleName*, a formula name *invName*, a table *symRules* storing rules, a table *symInvs* storing invariant formulas, a table *symCausalTab* storing causal relation

Output: An Isabelle proof script for a lemma: *lemmaWithProof*

```

1 asm1  $\leftarrow$  asmGenOnInv(symInvs, invName);
2 asm2  $\leftarrow$  asmGenOnRule(symRules, ruleName);
3 obtain1  $\leftarrow$  obtainGenOnInv(symInvs, invName);
4 obtain2  $\leftarrow$  obtainGenOnRule(symRules, ruleName);
5 symRelItem  $\leftarrow$ 
  | tbl_element(symCausalTab, (ruleName $\wedge$ invName));
6 proof  $\leftarrow$  caseAnalzI(symRelItem);
7 lemmaWithProof  $\leftarrow$  sprintf
8   | "lemma %sVs%s:
9   | assumes %s and %s
10  | shows invHoldForRule s f r (invariants N)
11  | proof - %s %s %s
12  | qed"
13  | (ruleName, invName, asm1, asm2, obtain1, obtain2, proof)
14 return lemmaWithProof

```

Due to length limitation, we illustrate the algorithm for generating a key part of the proof of the lemma *critVsinvl*: the generation of a subproof (e.g., lines 7-8) according to a symbolic relation tag of *invHoldRule₁₋₃*, which is shown in Algorithm 10. Input *relTag* is the result of the generalization step, which is discussed in Section V. In the body of function *rel2proof*, *sprintf* writes a formatted data to string and returns it. In line 10, *getFormField*(*relTag*) returns the field of formula *f'* if *relTag* = *invHoldRule₃*(*f'*). *rel2proof* transforms a symbolic relation tag into a paragraph of proof, as shown in lines 7-8, 10-11, or 13-14. If the tag is among *invHoldRule₁₋₂*, the transformation is rather straight-forward, else the form *f'* is assigned by the formula *getFormField*(*relTag*), and provided to tell Isabelle the formula which is used to construct the *invHoldRule₃* relation.

Algorithm 10: Generating a kind of proof which is according with a relation tag of *invHoldRule₁₋₃* : *rel2proof*

Input: A symbolic causal relation item *relTag*

Output: An Isabelle proof: *proof*

```

1 if relTag = invHoldRule1 then
2   | proof  $\leftarrow$  sprintf
3   | "have invHoldRule1 f r (invariants N)
4   | by(cut_tac a1 a2 b1, simp, auto)
5   | then have invHoldRule f r (invariants N) by blast";
6 else if relTag = invHoldRule2 then
7   | proof  $\leftarrow$  sprintf
8   | "have invHoldRule2 f r (invariants N) by(cut_tac a1 a2
9   | b1, simp, auto)
10  | then have invHoldRule f r (invariants N) by blast";
11 else
12   | f'  $\leftarrow$  getFormField(relTag);
13   | proof  $\leftarrow$  sprintf
14   | "have invHoldRule3 f r (invariants N)
15   | proof(cut_tac a1 a2 b1, simp, rule_tac x=%s in
16   | exI,auto)qed
17   | then have invHoldRule f r (invariants N) by blast"
18   | (symbf2Isabelle f');
19 return proof

```

VII. EXPERIMENTS

We implement our tool in Ocaml. Experiments are done with typical bus-snoopy benchmarks such as MESI and MOESI, as well as directory-based benchmarks such as German and FLASH. The detailed codes and experiment data can be found in [12]. Each experiment data includes the paraVerifier instance, invariant sets, Isabelle proof scripts. Experiment results are summarized in Table II.

Among all the work in the field of parameterized verification, only four of them have verified FLASH. The first full verification of safety properties of FLASH is done in [10]. Park and Dill proved the safety properties of FLASH using PVS. The CMP method, which adopts parameter abstraction and guard strengthening, is applied in [6] for verifying safety properties of FLASH. McMillan applied compositional model checking [11] and used Candence SMV to the verification of both safety and liveness properties of FLASH. Sylvain et.al have applied Cubeic to the verification FLASH [9], [13], which is theoretically based on an SMT model checking to the verification of array-based system. In the former three methods [10], [6], [11], auxiliary invariants are provided manually depending on verifier's deep insight in the FLASH protocol itself, while in Cubeic, auxiliary invariants are found automatically. In Cubeic, auxiliary invariants are searched backward by a heuristics-guided algorithm with the help of an oracle (a reference instance of the protocol), but these auxiliary invariants are in concrete form, and are not generalized to the parameterized form. Thus there is no parameterized proof derived for parameterized verification of FLASH.

The invariants-searching algorithm used in our work differs from that in Cubeic [9], [13] in that the heuristics in our work are based on the construction of causal relation which is uniquely proposed in our work. Thus the auxiliary invariants in

TABLE II
VERIFICATION RESULTS ON BENCHMARKS.

Protocols	#rules	#invariants	time (seconds)	Memory (MB)
mutualEx	4	5	3.25	7.3
MESI	4	3	2.47	11.5
MOESI	5	3	2.49	23.2
Germanish [9]	6	3	2.9	7.8
German [6]	13	52	38.67	14
FLASH_nodata	60	152	280	26
FLASH_data	62	162	510	26

our work are different from those found in [9], [13]. Moreover, we generalize these concrete invariants and causal relations into a parameterized proof, and generate a parameterized proof in Isabelle. The found invariants have abundant semantics reflecting the deep insight of the FLASH protocol design, and the readable Isabelle proof script formally proves these invariants. In this way, we prove the protocol with the highest assurance. To the best of knowledge, this work for the first time automatically generates a proof of safety properties of full version of FLASH in a theorem prover without auxiliary invariants manually provided by people.

VIII. CONCLUSION

The originality of `paraVerifier` lies in the following aspects: (1) instead of directly proving the invariants of a protocol by induction, we propose a general proof method based on the consistency lemma to decompose the proof goal into a number of small ones; (2) instead of proving the decomposed subgoals by hand, we automatically generate proofs for them based on the information of causal relation computed in a small protocol instance.

As we demonstrate in this work, combining theorem proving with automatic proof generation is promising in the field of formal verification of industrial protocols. Theorem proving can guarantee the rigorousness of the verification results, while automatic proof generation can release the burden of human interaction.

REFERENCES

- [1] Pnueli, A., Shahar, E.: A platform for combining deductive with algorithmic verification. In Alur, R., Henzinger, T., eds.: *Computer Aided Verification*. Volume 1102 of *Lecture Notes in Computer Science*. Springer Berlin Heidelberg (1996) 184–195
- [2] Björner, N., Browne, A., Manna, Z.: Automatic generation of invariants and intermediate assertions. *Theoretical Computer Science* **173**(1) (1997) 49 – 87
- [3] Arons, T., Pnueli, A., Ruah, S., Xu, Y., Zuck, L.: Parameterized verification with automatically computed inductive assertions? In: *Proc. 13th International Conference on Computer Aided Verification*. Volume 2102 of *LNCS*. Springer (2001) 221–234
- [4] Pnueli, A., Ruah, S., Zuck, L.: Automatic deductive verification with invisible invariants. In: *7th International Conference on Tools and Algorithms for the Construction and Analysis of Systems*. Volume 2031 of *LNCS*. Springer (2001) 82–97
- [5] Tiwari, A., Rueß, H., Saïdi, H., Shankar, N.: A technique for invariant generation. In: *Proc. 7th International Conference on Tools and Algorithms for the Construction and Analysis of Systems*. Volume 2031 of *LNCS*. Springer (2001) 113–127
- [6] Chou, C.T., Mannava, P., Park, S.: A simple method for parameterized verification of cache coherence protocols. In: *Proc. 5th International Conference on Formal Methods in Computer-Aided Design*. Volume 3312 of *LNCS*. Springer (2004) 382–398

- [7] Pandav, S., Slind, K., Gopalakrishnan, G.: Counterexample guided invariant discovery for parameterized cache coherence verification. In: *Proc. 13th IFIP Advanced Research Working Conference on Correct Hardware Design and Verification Methods*. Volume 3725 of *LNCS*. Springer (2005) 317–331
- [8] Lv, Y., Lin, H., Pan, H.: Computing invariants for parameter abstraction. In: *Proc. the 5th IEEE/ACM International Conference on Formal Methods and Models for Codesign, IEEE CS (2007)* 29–38
- [9] Conchon, S., Goel, A., Krstic, S., Mebsout, A., Zaïdi, F.: Cubicle: A parallel smt-based model checker for parameterized systems. In Madhusudan, P., Seshia, S., eds.: *Proc. 24th International Conference on Computer Aided Verification*. Volume 7358 of *LNCS*. Springer (2012) 718–724
- [10] Park, S., Dill, D.L.: Verification of flash cache coherence protocol by aggregation of distributed transactions. In: *Proceedings of the eighth annual ACM symposium on Parallel algorithms and architectures*, Padua, Italy, ACM (1996) 288–296
- [11] Mcmillan, K.L., Labs, C.B.: Parameterized verification of the flash cache coherence protocol by compositional model checking. In: *CHARME 01: IFIP Working Conference on Correct Hardware Design and Verification Methods, Lecture Notes in Computer Science 2144*, Springer (2001) 179–195
- [12] Y. Li, K.d.: `paraverifier` (2016) <https://github.com/paraVerifier/paraVerifier>.
- [13] Conchon, S., Goel, A., Krstic, S., Mebsout, A., Zaidi, F.: Invariants for finite instances and beyond. In: *FMCAD, Portland, Oregon, USA (October 2013)*